

# Determining the Sentiment of Financial News

The news database here will train the Naive Bayes, then RandomForest For deploying I'd recommend using the NewsAPI code shared and tag the sentiment via the trained NB.

## Constructing a Naive Bayes Classifier

- Load dataset
- Vectorize data
- Split data (80/20, train test, random\_state=0 so as to allow reproducibility)
- Initialize the NB classifier and fit
- Predict and measure accuracy

```
In [9]: 1 import numpy as np # linear algebra
2 import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
3 from sklearn.feature_extraction.text import CountVectorizer
4
5 news_pd = pd.read_csv("./news_with_sentiment.csv")
6 news_pd = news_pd[:2000] # 28,000 rows will use more RAM than is available
7
8 cv = CountVectorizer() # Convert text data to a vector as that is required
9 X = cv.fit_transform(news_pd['text']).toarray()
10 y = news_pd['sentiment'] # y = the variable we are trying to predict
```

```
In [10]: 1 # Split train and test data (80/20)
2 from sklearn.model_selection import train_test_split
3 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
4
5 # Initialize the Gaussian Naive Bayes Classifier, then fit the data
6 from sklearn.naive_bayes import GaussianNB
7 classifier = GaussianNB()
8 classifier.fit(X_train, y_train)
```

```
Out[10]: GaussianNB(priors=None)
```

```
In [11]: 1 # Predict sentiment of our test data
2 y_pred = classifier.predict(X_test)
3
4 from sklearn.metrics import accuracy_score
5 score = accuracy_score(y_test, y_pred)
```

And now we can view the accuracy:

0.685

1 Roughly 68% accuracy. Not exactly stellar, if you reduce the dataset further you end up with higher accuracy which is interesting.

In [14]:

```
1 news_pd = pd.read_csv("./news_with_sentiment.csv")
2 news_pd = news_pd[:1000] # 28,000 rows will use more RAM than is av
3
4 cv = CountVectorizer()
5 X = cv.fit_transform(news_pd['text']).toarray()
6 y = news_pd['sentiment']
7
8 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size
9
10 classifier = GaussianNB()
11 classifier.fit(X_train, y_train)
12
13 y_pred = classifier.predict(X_test)
14
15 score = accuracy_score(y_test, y_pred)
16
17 print(score)
```

0.775

77.5% accuracy on a 1000 row dataset with an 80/20 split.

After research, Naive Bayes appears to be better with smaller datasets but perhaps we can improve:

## To improve on our Naive Bayes we can now try a Random Forest:

- Load dataset
- Remove stopwords, min\_df=7 means the data is irrelevant if used in more than 7 documents, max\_df of 0.8 means it also is irrelevant if used in more than 80% of documents
- Vectorize data (max\_features is the max number of WORDS in Vector form that will influence the sentiment)
- Split data (80/20, train test, random\_state=0 so as to allow reproducibility)
- Initialize the Random Forest classifier and fit
- Predict and measure accuracy

Azure Notebooks (/#) Preview (https://docs.microsoft.com/en-us/azure/notebooks/) My Projects (https://docs.microsoft.com/en-us/azure/notebooks/)

```
In [15]: 1 # Read in 20,000 headlines (https://data.1302/projects#)
2 news_pd = pd.read_csv("./news_with_sentiment.csv")
3 news_pd = news_pd[:20000] # 28,000 rows will use more RAM than is a
4 y = news_pd['sentiment']
```

```
In [16]: 1 from nltk.corpus import stopwords
2 from sklearn.feature_extraction.text import TfidfVectorizer
3
4 # Remove stopwords and vectorize the dataset
5 #TfidfVectorizer converts a collection of raw documents to a matrix
6 vectorizer = TfidfVectorizer(max_features=2500, min_df=7, max_df=0.
7 processed_features = vectorizer.fit_transform(news_pd['text']).toar
```

```
In [17]: 1 # 80/20 data split
2 from sklearn.model_selection import train_test_split
3 X_train, X_test, y_train, y_test = train_test_split(processed_featu
4
5 # Fit our model with split data, starting with 450 estimators (450
6 from sklearn.ensemble import RandomForestClassifier
7
8 text_classifier = RandomForestClassifier(n_estimators=450, random_s
9 text_classifier.fit(X_train, y_train)
```

/home/nbuser/anaconda3\_420/lib/python3.5/site-packages/sklearn/ensembl  
e/weight\_boosting.py:29: DeprecationWarning: numpy.core.umath\_tests is  
an internal NumPy module and should not be imported. It will be remove  
d in a future NumPy release.

from numpy.core.umath\_tests import inner1d

```
Out[17]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='g
ini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=450, n_jobs=1,
                                oob_score=False, random_state=0, verbose=0, warm_start=False)
```

```
In [18]: 1 # Predicting the sentiment of our test data
2 predictions = text_classifier.predict(X_test)
3
4
5 # Checking our accuracy
6 from sklearn.metrics import accuracy_score
7 print(accuracy_score(y_test, predictions))
```

0.93575

93.57% accuracy

## Hyperparameter Tuning:

- Choose a set of trees we want to test
- Train the model with n trees. store accuracv


- Loop above until complete
- Plot the resulting trees v accuracy

In [ ]:

```

1  from sklearn.ensemble import RandomForestRegressor
2
3  estimators = [50, 100, 150, 200, 250, 300, 350, 400, 450, 500, 550,
4  accuracy = []
5
6  for estimator_num in estimators:
7      # Fit and predict
8      text_classifier = RandomForestClassifier(n_estimators=estimator_num)
9      text_classifier.fit(X_train, y_train)
10     predictions = text_classifier.predict(X_test)
11
12     # Store accuracy
13     from sklearn.metrics import accuracy_score
14     accuracy.append(accuracy_score(y_test, predictions))
15
16
17 # Graph reported accuracy of various sets of estimators
18 import matplotlib.pyplot as plt
19
20 plt.plot(estimators, accuracy)
21 plt.ylabel('Accuracy')
22 plt.xlabel('Estimators')
23 plt.show()
24
25 print(estimators)
26 print(accuracy)

```



A strange curve?

As per: [https://en.wikipedia.org/wiki/Talk%3ARandom\\_forest](https://en.wikipedia.org/wiki/Talk%3ARandom_forest)  
[https://en.wikipedia.org/wiki/Talk%3ARandom\\_forest](https://en.wikipedia.org/wiki/Talk%3ARandom_forest)

"Random Forests does not overfit. The testing performance of Random Forests does not decrease (due to overfitting) as the number of trees increases. Hence after certain number of trees the performance tend to stay in a certain value."

Azure  
Notebooks  
(/#)

Preview  
(/help/preview)

My  
Projects  
(g00311302/projects#)

Help  
(https://docs.microsoft.com/en-us/azure/notebooks/)

However, we can also see that ~250 estimators/trees is the ideal parameter.

Naive Bayes v Random Forest v SVM:

[https://www.researchgate.net/publication/336225950\\_Comparison\\_of\\_Naive\\_Bayes\\_Supervised\\_Learning\\_Algorithms](https://www.researchgate.net/publication/336225950_Comparison_of_Naive_Bayes_Supervised_Learning_Algorithms)  
([https://www.researchgate.net/publication/336225950\\_Comparison\\_of\\_Naive\\_Bayes\\_Supervised\\_Learning\\_Algorithms](https://www.researchgate.net/publication/336225950_Comparison_of_Naive_Bayes_Supervised_Learning_Algorithms))

## Pull Fresh News:

Azure In [19]: Preview My Projects Help (https://docs.microsoft.com/en-us/azure/notebooks/)

```

1 import requests
2 import time
3 import datetime
4
5
6 articleCount = 0
7
8 headers = {
9     'User-Agent': 'Mozilla/5.0 (Windows NT 6.1) AppleWebKit/537.36
10 }
11
12 stocks = ['TSLA', 'AMZN', 'MMM', 'INTC', 'GOOGL', 'FB', 'MSFT', 'AA
13 list_of_headlines = []
14 for line in stocks:
15     ticker = line
16
17     try:
18
19         #Query for the stock name, for refined news queries.
20         resp = requests.get(
21             url="https://www.alphavantage.co/query?function=SYMBOL_
22                 ticker), headers=headers)
23         data = resp.json()
24         companyName = data['bestMatches'][0]['2. name']
25         print("Company Name: " + companyName)
26
27         #Query for news
28         resp = requests.get(
29             url='https://newsapi.org/v2/everything?'
30             'q={}&'
31             'from=2020-01-05' # This is the OLDEST date an article can be from,
32             'sortBy=popularity&' #Filter by popularity (read the newsapi docs)
33             'apiKey=fe00115ceffe418988616191b03e1c74'.format(
34                 ticker + " " + companyName), headers=headers) #Add
35         data = resp.json()
36
37         for article in data['articles']:
38             articleCount = articleCount + 1
39             newsTitle = article['title']
40             print(newsTitle)
41             list_of_headlines.append(newsTitle)
42
43         time.sleep(1)
44
45     except Exception as e:
46         print("Error: " + str(e))
47         time.sleep(10)
48
49 # Create the pandas DataFrame and save to csv
50 df = pd.DataFrame({'headlines':list_of_headlines})
51 df.to_csv('fresh_news_month_tsla.csv', encoding='utf-8', mode='w',

```

Cramer Weighs In On Cracker Barrel, UPS And More

AWS Announces General Availability of Amazon Keyspaces (for Apache C  
assandra)

Company Name: 3M Company

Dow Jones 378-Point Intraday Gain Fades, But 3M A Bright Spot; Netfl  
ix, Tesla Weigh In, Nvidia, Twitter's Business Daily

Azure  
Notebooks  
(/#)

ix, Tesla Weigh In Nasdaq - Investor's Business Daily  
Dow Jones, US Stocks Rise As Countries Begin To Reopen Economies - Investor's Business Daily  
3M's stock surges on earnings beat, that was nearly 20 years in the making  
3M Co (MMM) Q1 2020 Earnings Call Transcript  
Is 3M Oversold At \$155?  
3M Company (MMM) CEO Mike Roman on Q1 2020 Results - Earnings Call Transcript  
Were Hedge Funds Right About 3M Company (MMM)?  
Did You Acquire 3M (MMM) Before February 9, 2017? Johnson Fintel Continues its Investigation of 3M; Should Management be Held Accountable for Investors Losses?  
3M (MMM) Gains But Lags Market: What You Should Know  
3M Holds Good On Its Promise To Prioritize Dividend

Azure In [20]: Preview My Projects (https://docs.microsoft.com/en-us/azure/notebooks/) (/#)

```

1 from nltk.corpus import stopwords
2 from sklearn.feature_extraction.text import TfidfVectorizer
3 from sklearn.feature_extraction.text import CountVectorizer
4
5 # Read in fresh news
6 fresh_news = pd.read_csv('./fresh_news_month_tsla.csv')
7 fresh_news['headlines'].head(5)
8
9 # Vectorize new text data with a max of 40 words being predictors
10 vectorizer_new_data = CountVectorizer(max_features=40, min_df=9)
11 processed_features_new_data = vectorizer_new_data.fit_transform(fresh_news['headlines'])
12
13 # Vectorize training text data with a max of 40 words being predictors
14 vectorizer = CountVectorizer(max_features=40, min_df=9)
15 processed_features = vectorizer.fit_transform(news_pd['text']).toarray()
16
17 X_train, X_test, y_train, y_test = train_test_split(processed_features, y_train,
18                                                     test_size=0.2, random_state=42)
19
20 # Predict on new/fresh news after fitting on training data
21 text_classifier = RandomForestClassifier(n_estimators=650, random_state=42)
22 text_classifier.fit(X_train, y_train)
23
24 predictions = text_classifier.predict(processed_features_new_data)
25
26 # Output our predictions
27 print(predictions)
28
29 for i in range(len(predictions)):
30     if predictions[i] == 1:
31         print("Positive: " + fresh_news['headlines'][i])
32     if predictions[i] == -1:
33         print("Negative: " + fresh_news['headlines'][i])

```

```

[ 0  0  0  0  1  1  0  0 -1  0  0  0  0  0  0  0  1  0  1  1  1 -1
 0  0
 0  1  0 -1  0  0  0  0  0  0  0  0  0  0  0  0  0  0  1  0  0  0
 0  1
 1  1  1  0  1  0  0  0 -1  0  0 -1  0  0  0  0  0 -1  0  0  0  0
 0  0
 0  0  0 -1  0  0  0  0  1  0 -1  0  0  0  0 -1  1  0  1  0  0  0
 0  1
 0  0  0  0  1  0  0  0  0  0  0  0  1  0  0  0  0  0  0  1 -1  0
 0  0
 0  1  0  0 -1  0  0  0  0  0  0  0 -1  1  1  0  0  0  1  0  0  1
 0  0
 0  0  1  0  0  0  0  0  1  0  0  0  0  0  1  0]
Positive: Tesla wants to reopen California factory, but local author
ities say not yet
Positive: Ford is first auto maker to warn of lower sales, but unlik
ely to be last
Negative: Market Extra: The S&P 500 just posted the most daily swing
s of 3% or greater in more than a decade—even as the stock market hi
ts a 5-week high
Positive: The force that’s propelled the stock market rally will exh
aust itself this week
Positive: Tesla Confirms Shanghai Gigafactory Shutdown, But Says I

```



It's All 'According To Plan'

Positive: Steve Grasso Says Tesla Has Defied All Laws Of Probability  
y' Projects

Positive: Amazon stock hits record high on hopes for a coronavirus-related boom

Negative: AMC's stock soars after report Amazon held merger talks

Positive: Netflix stock surges to record high as investors bet on streaming during coronavirus

Negative: Market Extra: The S&P 500 just posted the most daily swings of 3% or greater in more than a decade—even as the stock market hits a 5-week high

Positive: Dow Jones, US Stocks Rise As Countries Begin To Reopen Economies - Investor's Business Daily

Positive: Did You Acquire 3M (MMM) Before February 9, 2017? Johnson  
Fistel Continues its Investigation of 3M; Should Management be Held  
Accountable for Investors Losses?

## Positive: 3M (MMM) Gains But Lags Market: What You Should Know

## Positive: 3M Holds Good On Its Promise To Prioritize Dividend

Positive: Stocks fights for gains as earnings season revs up - Fox B  
usiness

## Positive: The Importance Of Reading Footnotes - Uncovering Material Items In Filings

## Negative: Dow Blue-Chip 3M Co Surges on Coronavirus Demand, but Bond King Warns of Danger

## Negative: Why I Think You Should Buy This Defence Stock Before May

## Negative: Call Traders Blast These 2 Chip Stocks

Negative: 5 Tech Stocks Poised to Beat Estimates This Earnings Season - Yahoo Finance

Positive: Making Most Of Lockdowns, Facebook Gaming Launches Earlier Than Planned - Benzinga

Negative: Alphabet Announces First Quarter 2020 Results (Alphabet)

Negative: Market Extra: The S&P 500 just posted the most daily swings of 3% or greater in more than a decade—even as the stock market hits a 5-week high

Positive: Big Data tech CEO on the federal government's response to coronavirus: 'A total failure of leadership'

Positive: Dropbox's first quarterly profit is a sign of the ever-changing economy

Positive: Investors have \$5.1 trillion hiding out in the shares of five companies, which will be tested this week

Positive: Making Most Of Lockdowns, Facebook Gaming Launches Earlier Than Planned - Benzinga

Positive: Dropbox's first quarterly profit is a sign of the ever-changing economy

Positive: Investors have \$5.1 trillion hiding out in the shares of five companies, which will be tested this week

Negative: FB to allow employees to work remotely until year end

## Positive: How Large Option Traders Are Playing Microsoft As Cloud Business Booms - Yahoo Finance

Negative: Tech Stocks' Apr 29 Earnings Lineup: NOW, FB, MSFT, FICO, GIB - Yahoo Finance

## Negative: Hedge Funds' #3 Stock Pick Debunked Naysayers

## Positive: Optimism May Be Over Done In The Equity Markets

## Positive: Jonathan Angrist's Cognios Can't Deliver Despite Apple, Amazon, Microsoft Bets

Positive: Microsoft announces registered exchange offers

Positive: The Ratings Game: Why Apple investors should be worried by

Azure  
Notebooks  
(/#)

Preview (https://docs.microsoft.com/en-us/azure/notebooks/) My Projects (/00311302/projects#) Help (https://docs.microsoft.com/en-us/azure/notebooks/)

AT&T's earnings  
Positive: Big Data tech CEO on the federal government's response to coronavirus: 'A total failure of leadership'  
Positive: Investors have \$5.1 trillion hiding out in the shares of five companies, which will be tested this week  
Positive: Gene Munster Dismisses Goldman's Apple Downgrade, Says Cupertino Has Long-Term Earnings Power

Save our model to disk for production deployment to Sparkbot

```
In [ ]: 1 from joblib import dump
        2 dump(text_classifier, 'sentimentclassified.joblib')
```

```
In [ ]: 1
```